

Spring 2025: CS5720

Neural Networks and Deep Learning - ICP-5

Gonaboyina Vijay Vardhan (700755141)

Github link: <https://github.com/Vijayvardhan02/NEURAL-NETWORK-DEEP-LEARNING-ICP5>

Video link:

[https://drive.google.com/file/d/1accMtlVqds5gahCCb5RBXJyU6qKDFhE1/view?usp=drive link](https://drive.google.com/file/d/1accMtlVqds5gahCCb5RBXJyU6qKDFhE1/view?usp=drive_link)

- 1. Follow the instruction below and then report how the performance changed. (apply all at once)**
 - Convolutional input layer, 32 feature maps with a size of 3×3 and a rectifier activation function.
 - Dropout layer at 20%.
 - Convolutional layer, 32 feature maps with a size of 3×3 and a rectifier activation function.
 - Max Pool layer with size 2×2 .
 - Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function.
 - Dropout layer at 20%.
 - Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function.
 - Max Pool layer with size 2×2 .
 - Convolutional layer, 128 feature maps with a size of 3×3 and a rectifier activation function.
 - Dropout layer at 20%.
 - Convolutional layer, 128 feature maps with a size of 3×3 and a rectifier activation function.
 - Max Pool layer with size 2×2 .
 - Flatten layer.
 - Dropout layer at 20%.
 - Fully connected layer with 1024 units and a rectifier activation function.
 - Dropout layer at 20%.

- Fully connected layer with 512 units and a rectifier activation function.
 - Dropout layer at 20%.
 - Fully connected output layer with 10 units and a Softmax activation function
- Did the performance change?
2. **Predict the first 4 images of the test data using the above model. Then, compare with the actual label for those 4 images to check whether or not the model has predicted correctly.**
 3. **Visualize Loss and Accuracy using the history object**

Code:

```
✓ 1m ▶ import tensorflow as tf
from tensorflow.keras import layers, models
import matplotlib.pyplot as plt

# Load the dataset (e.g., CIFAR-10)
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar10.load_data()

# Normalize pixel values to between 0 and 1
x_train, x_test = x_train / 255.0, x_test / 255.0

# Create the model
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
    layers.Dropout(0.2),
    layers.Conv2D(32, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),

    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.Dropout(0.2),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),

    layers.Conv2D(128, (3, 3), activation='relu'),
    layers.Dropout(0.2),
    layers.Conv2D(128, (3, 3), activation='relu'),
    # Remove one max pooling layer here to prevent dimension collapse
    layers.Flatten(),
    layers.Dropout(0.2),
    layers.Dense(1024, activation='relu'),
    layers.Dropout(0.2),
```



```
layers.Dense(512, activation='relu'),
layers.Dropout(0.2),
layers.Dense(10, activation='softmax')
])

# Compile the model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(x_train, y_train, epochs=10, validation_data=(x_test, y_test))

# Evaluate the model on the test data
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)

print(f"Test accuracy: {test_acc}")
print(f"Test loss: {test_loss}")

# Predict the first 4 images
predictions = model.predict(x_test[:4])
for i, prediction in enumerate(predictions):
    print(f"Predicted label for image {i+1}: {prediction.argmax()} (Actual label: {y_test[i][0]})")

# Visualizing the loss and accuracy using the history object
plt.figure(figsize=(12, 6))
# Plotting loss
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Loss during training')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

# Plotting accuracy
plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Accuracy during training')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```

Output:

```
/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape` to `input_dim`  
super().__init__(activity_regularizer=activity_regularizer, **kwargs)  
Epoch 1/10  
1563/1563 — 208s 130ms/step - accuracy: 0.2766 - loss: 1.9018 - val_accuracy: 0.4703 - val_loss: 1.4382  
Epoch 2/10  
1563/1563 — 261s 130ms/step - accuracy: 0.4907 - loss: 1.4007 - val_accuracy: 0.5250 - val_loss: 1.3330  
Epoch 3/10  
1563/1563 — 208s 133ms/step - accuracy: 0.5668 - loss: 1.2128 - val_accuracy: 0.6161 - val_loss: 1.0945  
Epoch 4/10  
1563/1563 — 260s 132ms/step - accuracy: 0.6134 - loss: 1.1116 - val_accuracy: 0.6393 - val_loss: 1.0429  
Epoch 5/10  
1563/1563 — 259s 130ms/step - accuracy: 0.6402 - loss: 1.0292 - val_accuracy: 0.6459 - val_loss: 0.9990  
Epoch 6/10  
1563/1563 — 265s 132ms/step - accuracy: 0.6648 - loss: 0.9637 - val_accuracy: 0.6754 - val_loss: 0.9467  
Epoch 7/10  
1563/1563 — 201s 128ms/step - accuracy: 0.6850 - loss: 0.9187 - val_accuracy: 0.6858 - val_loss: 0.9070  
Epoch 8/10  
1563/1563 — 206s 131ms/step - accuracy: 0.6948 - loss: 0.8886 - val_accuracy: 0.7090 - val_loss: 0.8535  
Epoch 9/10  
1563/1563 — 199s 127ms/step - accuracy: 0.7054 - loss: 0.8579 - val_accuracy: 0.7075 - val_loss: 0.8560  
Epoch 10/10  
1563/1563 — 205s 129ms/step - accuracy: 0.7176 - loss: 0.8224 - val_accuracy: 0.7132 - val_loss: 0.8441  
313/313 - 9s - 29ms/step - accuracy: 0.7132 - loss: 0.8441  
Test accuracy: 0.7131999731063843  
Test loss: 0.8440958857536316  
1/1 — 0s 163ms/step  
Predicted label for image 1: 3 (Actual label: 3)  
Predicted label for image 2: 8 (Actual label: 8)  
Predicted label for image 3: 8 (Actual label: 8)  
Predicted label for image 4: 0 (Actual label: 0)
```

